

Evaluation Metrics Report - Data Vortex A'26

Team: Event Horizon | Round 2: Sentiment Classification

Page Map

Page	Section
1	Page Map (this page)
2	Overall Performance Metrics (Accuracy, Macro-F1, Weighted-F1, Precision, Recall)
2	Per-Class Breakdown Table (Negative / Neutral / Positive)
2	BERTweet Benchmark Comparison (if available)
3	Confusion Matrix Visualization
3	Per-Class Metrics Visualization

Evaluation Metrics Report - Data Vortex A'26

Team: Event Horizon | Round 2: Sentiment Classification

1. Overall Performance Metrics

Metric	Value
Model	Linear SVM
Test Accuracy	0.5911 (baseline: 0.3481 lift: +0.2430)
Macro-F1	0.5914 (baseline: 0.1721 lift: +0.4193)
Weighted-F1	0.5892
Macro Precision	0.5919
Macro Recall	0.5954

2. Per-Class Breakdown

Class	Precision	Recall	F1-Score	Support
Negative	0.5930	0.7016	0.6428	486
Neutral	0.5283	0.4927	0.5099	550
Positive	0.6545	0.5919	0.6216	544
macro avg	0.5919	0.5954	0.5914	1580
weighted avg	0.5916	0.5911	0.5892	1580

3. BERTweet Benchmark Comparison [PRE-COMPUTED, GPU-TRAINED]

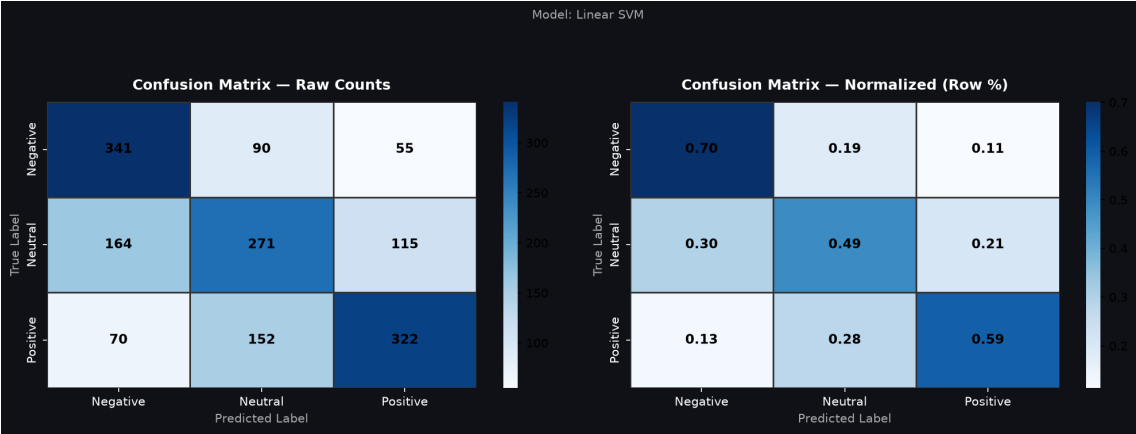
BERTweet was fine-tuned on Google Colab T4 GPU (not locally reproduced). Verify metrics independently: `python benchmarks/bertweet/verify_benchmark.py`

Metric	Linear SVM (Local)	BERTweet (Colab GPU)
Model	Linear SVM	vinai/bertweet-base
Macro-F1	0.5914	0.7330
Accuracy	0.5911	0.7335
Kappa	N/A	0.6015
Split	80/20 train/test	70_10_20_clean_dedup

Evaluation Metrics Report - Data Vortex A'26

Team: Event Horizon | Round 2: Sentiment Classification

4. Confusion Matrix (Visualized)



5. Per-Class Metrics (Visualized)

