

Evaluation Metrics Report - Data Vortex A'26

Team: Event Horizon | Round 2: Sentiment Classification

Page Map

Page	Section
1	Page Map (this page)
2	Overall Performance Metrics (Accuracy, Macro-F1, Weighted-F1, Precision, Recall)
2	Per-Class Breakdown Table (Negative / Neutral / Positive)
2	BERTweet Benchmark Comparison (if available)
3	Confusion Matrix Visualization
3	Per-Class Metrics Visualization

Evaluation Metrics Report - Data Vortex A'26

Team: Event Horizon | Round 2: Sentiment Classification

1. Overall Performance Metrics

Metric	Local Model	BERTweet (Colab GPU)
Model	Logistic Regression	vinai/bertweet-base
Test Accuracy	0.5994 (baseline: 0.3478)	0.7335
Macro-F1	0.6005 (baseline: 0.1720)	0.7330
Weighted-F1	0.5989	N/A
Macro Precision	0.6028	N/A
Macro Recall	0.6029	N/A
Cohen's Kappa	N/A	0.6015

2. Per-Class Breakdown

Class	Precision	Recall	F1-Score	Support
Negative	0.5915	0.6914	0.6376	486
Neutral	0.5362	0.5255	0.5308	550
Positive	0.6808	0.5919	0.6332	544
macro avg	0.6028	0.6029	0.6005	1580
weighted avg	0.6030	0.5994	0.5989	1580

3. BERTweet Benchmark Comparison [PRE-COMPUTED, GPU-TRAINED]


BERTweet was fine-tuned on Google Colab T4 GPU (not locally reproduced). Verify metrics independently: `python benchmarks/bertweet/verify_benchmark.py`

Metric	Local Model (Logistic Regression)	BERTweet (Colab GPU)
Model	Logistic Regression	vinai/bertweet-base
Macro-F1	0.6005	0.7330
Accuracy	0.5994	0.7335
Kappa	N/A	0.6015
Split	80/20 train/test	70_10_20_clean_dedup

Evaluation Metrics Report - Data Vortex A'26

Team: Event Horizon | Round 2: Sentiment Classification

4. Confusion Matrix (Visualized)



5. Per-Class Metrics (Visualized)

